

ISL Fall 2015

Assignment II. 100 pts.

NAME(s) :

LOKESHWAR REDDY INJA (16207239/ LIRK9)

SANTHOSH KUMAR GATTU (16211118/ SG6N6)

SANTOSH KUMAR (16207971/ SK7Z9)

You may submit this assignment in groups of upto three each. Write your names on this sheet and include it as the cover page for your submission.

The objective of this assignment is to practice using R, and gain a fundamental understanding of linear regression. Your submission should include both your code as well your answers to the questions.

Electronic submission on Blackboard is due **latest by 11 pm on Wed, Sep 30th**.

You may upload upto **three** submissions **before** the deadline - only the last submission will be graded. Submissions received after the deadline will be graded only for effort for a maximum of 70% of the total grade (Refer to class syllabus for detailed grading policy).

State any assumptions you make, justify your answers, show intermediate steps and explain your results for maximum credit. All answers should be in your own words with any sources you refer to cited at the appropriate places. Any knowledge you acquire from the Internet should be written in your own words and be appropriately referenced. Copying and pasting from the Internet, each other or any other source will not count as your effort (Refer to class syllabus for detailed policy on plagiarism).

Remember that answers need to be word-processed (NOT handwritten) and should use R.

Answer the following questions from Chapter 3.

Undergraduate Students:

Questions 4, 8, 10

Graduate Students:

Questions 6, 10, 15

Question 6.

Using (3.4), argue that in the case of simple linear regression, the least squares line always passes through the point (\bar{x}, \bar{y}) .

Solution:

We know that the Simple linear regression is used to predict the value of response Y on the basis of predictor X. It also assumes the relationship between X and Y to be approximately linear.

In mathematical terms we can represent Y as a function of X.

$$Y = f(X) \cong \beta_0 + \beta_1 X \text{ -----} \rightarrow \text{equation 1}$$

Substitute (\bar{X}, \bar{Y}) in equation 1

$$\bar{Y} = \beta_0 + \beta_1 \bar{X}$$

$$\beta_0 = \bar{Y} - \beta_1 \bar{X} \text{ -----} \rightarrow \text{equation 2}$$

Now again substitute β_0 value in equation 1

$$\text{We get, } Y = \bar{Y} - \beta_1 \bar{X} + \beta_1 X \text{ ---} \rightarrow \text{equation 3}$$

Now from the equation 3 it is clear that the point (\bar{X}, \bar{Y}) will always pass through it.


If we pass the point through the equation, we get

$$\bar{Y} = \bar{Y} - \beta_1 \bar{X} + \beta_1 \bar{X}$$

$$0 = 0 \text{ -----} \rightarrow \text{equation 4}$$

$$0 = 0$$

Hence, with the help of above argument we can argue that in the case of simple linear regression, the least squares line will always pass through the point (\bar{X}, \bar{Y}) .



X

Question 10.

This question should be answered using the **Carseats** data set.

(a) Fit a multiple regression model to predict **Sales** using **Price**, **Urban**, and **US**.



(b)

Interpretation of each coefficient from the below data.

```
> summary(saleslm)
```

Call:

```
lm(formula = Sales ~ Price + Urban + US, data = Carseats)
```

Residuals:

Min	1Q	Median	3Q	Max
-6.9206	-1.6220	-0.0564	1.5786	7.0581

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	13.043469	0.651012	20.036	< 2e-16 ***
Price	-0.054459	0.005242	-10.389	< 2e-16 ***
UrbanYes	-0.021916	0.271650	-0.081	0.936
USYes	1.200573	0.259042	4.635	4.86e-06 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.472 on 396 degrees of freedom

Multiple R-squared: 0.2393, Adjusted R-squared: 0.2335

F-statistic: 41.52 on 3 and 396 DF, p-value: < 2.2e-16

- **Price:** On an average when price increases by 1 dollar the sales decreases by 54.459 units when all other predictors remain fixed because the coefficient of the predictor price is -0.054459. The standard error is also less in the model.
- **Urban:** the unit sales in urban areas are 21.916 units less when compared to the unit sales in rural areas when all other predictors remain fixed. Urban predictor cannot be used as it did not reject the null hypothesis which is identified from the large 'p-value' (0.936)
- **US :** the unit sales in US region are 1200.573 units greater than the sales in non-US regions when all other predictors remain fixed.

(C)

Model in Equation form:

$$\text{Sales} = B_0 + B_1 * (\text{price}) + B_2 * (\text{Urban}) + B_3 * (\text{US}) + \text{error}$$

Where $B_0 = 13.043469$, $B_1 = -0.054459$ $B_2 = -0.021916$ $B_3 = 1.200573$

Equation for Urban and non US is:

$$\text{Sales} = B_0 + B_1 * \text{price} + B_2 * \text{Urban} + \text{error}$$

Equation for rural and non US is:

$$\text{Sales} = B_0 + B_1 * \text{price} + \text{error}$$

Equation for Urban and US is:

$$\text{Sales} = B_0 + B_1 * \text{price} + B_2 * \text{Urban} + B_3 * \text{US} + \text{error}$$

Equation for rural and US is:

$$\text{Sales} = B_0 + B_1 * \text{price} + B_3 * \text{US} + \text{error}$$

(D)

Price predictor rejects the Null hypothesis as p-value ($2e-16$) is very less.

US predictor rejects the Null hypothesis as p-value ($2e-16$) is very less.

(E)

Please find below attached the .R file



(F)

Summary for the fit in (10 a) is

```
> summary(saleslm)
```

Call:

```
lm(formula = Sales ~ Price + Urban + US, data = Carseats)
```

Residuals:

Min	1Q	Median	3Q	Max
-6.9206	-1.6220	-0.0564	1.5786	7.0581

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	13.043469	0.651012	20.036	< 2e-16 ***
Price	-0.054459	0.005242	-10.389	< 2e-16 ***
UrbanYes	-0.021916	0.271650	-0.081	0.936
USYes	1.200573	0.259042	4.635	4.86e-06 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.472 on 396 degrees of freedom

Multiple R-squared: 0.2393, Adjusted R-squared: 0.2335

F-statistic: 41.52 on 3 and 396 DF, p-value: < 2.2e-16

Summary for (10.e) is

```
> summary(saleslm1)
```

Call:

```
lm(formula = Sales ~ Price + US, data = Carseats)
```

Residuals:

Min	1Q	Median	3Q	Max
-6.9269	-1.6286	-0.0574	1.5766	7.0515

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	13.03079	0.63098	20.652	< 2e-16 ***
Price	-0.05448	0.00523	-10.416	< 2e-16 ***
USYes	1.19964	0.25846	4.641	4.71e-06 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.469 on 397 degrees of freedom

Multiple R-squared: 0.2393, Adjusted R-squared: 0.2354

F-statistic: 62.43 on 2 and 397 DF, p-value: < 2.2e-16

Both the models fit the data very poorly as the R-squared value is 0.2393. In both models 10.e model is better than 10.a model because RSE is less for 10.e model.

(G)

Please find below attached the .R file



	2.5 %	97.5 %
(Intercept)	11.79032020	14.27126531
Price	-0.06475984	-0.04419543
USYes	0.69151957	1.70776632


(H)



Please find below attached the .R file




Standardized residuals versus leverage plots show the presence of a few outliers and some leverage points.

The plots for the model 10.e:







Question 15.

(a)

Expect 'chas' all the other predictors have less p-value, so all predictors are significant except chas because chas did not reject Null hypothesis. Also, chas has very low R-squared (0.003124) value, so it is not a good predictor.

'Rad' and 'tax' are very good predictor among all predictors because they have very high R-squared (for rad-0.3913 and for tax-0.3396) value.

Please find attached below the .R file.



isl15a.R

```
> lplot(crim,zn)
```

Call:

```
lm(formula = y ~ x, data = Boston)
```

Residuals:

Min	1Q	Median	3Q	Max
-4.429	-4.222	-2.620	1.250	84.523

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	4.45369	0.41722	10.675	< 2e-16 ***
x	-0.07393	0.01609	-4.594	5.51e-06 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 8.435 on 504 degrees of freedom
Multiple R-squared: 0.04019, Adjusted R-squared: 0.03828
F-statistic: 21.1 on 1 and 504 DF, p-value: 5.506e-06

```
> lmplot(crim,indus)
```

Call:
lm(formula = y ~ x, data = Boston)

Residuals:

Min	1Q	Median	3Q	Max
-11.972	-2.698	-0.736	0.712	81.813

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-2.06374	0.66723	-3.093	0.00209 **
x	0.50978	0.05102	9.991	< 2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 7.866 on 504 degrees of freedom
Multiple R-squared: 0.1653, Adjusted R-squared: 0.1637
F-statistic: 99.82 on 1 and 504 DF, p-value: < 2.2e-16

```
> lmplot(crim,chas)
```

Call:
lm(formula = y ~ x, data = Boston)

Residuals:

Min	1Q	Median	3Q	Max
-3.738	-3.661	-3.435	0.018	85.232

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	3.7444	0.3961	9.453	<2e-16 ***
x	-1.8928	1.5061	-1.257	0.209

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 8.597 on 504 degrees of freedom
Multiple R-squared: 0.003124, Adjusted R-squared: 0.001146
F-statistic: 1.579 on 1 and 504 DF, p-value: 0.2094

```
> lmplot(crim,nox)
```

Call:

```
lm(formula = y ~ x, data = Boston)
```

```
Residuals:
```

Min	1Q	Median	3Q	Max
-12.371	-2.738	-0.974	0.559	81.728

```
Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-13.720	1.699	-8.073	5.08e-15 ***
x	31.249	2.999	10.419	< 2e-16 ***

```
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 7.81 on 504 degrees of freedom
```

```
Multiple R-squared:  0.1772, Adjusted R-squared:  0.1756
```

```
F-statistic: 108.6 on 1 and 504 DF, p-value: < 2.2e-16
```

```
> lmplot(crim,rm)
```

```
Call:
```

```
lm(formula = y ~ x, data = Boston)
```

```
Residuals:
```

Min	1Q	Median	3Q	Max
-6.604	-3.952	-2.654	0.989	87.197

```
Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	20.482	3.365	6.088	2.27e-09 ***
x	-2.684	0.532	-5.045	6.35e-07 ***

```
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 8.401 on 504 degrees of freedom
```

```
Multiple R-squared:  0.04807, Adjusted R-squared:  0.04618
```

```
F-statistic: 25.45 on 1 and 504 DF, p-value: 6.347e-07
```

```
> lmplot(crim,age)
```

```
Call:
```

```
lm(formula = y ~ x, data = Boston)
```

```
Residuals:
```

Min	1Q	Median	3Q	Max
-6.789	-4.257	-1.230	1.527	82.849

```
Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-3.77791	0.94398	-4.002	7.22e-05 ***
x	0.10779	0.01274	8.463	2.85e-16 ***

```
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 8.057 on 504 degrees of freedom
```

```
Multiple R-squared:  0.1244, Adjusted R-squared:  0.1227
```

```
F-statistic: 71.62 on 1 and 504 DF, p-value: 2.855e-16
```

```
> lmplot(crim,dis)
```

```
Call:
```

```
lm(formula = y ~ x, data = Boston)
```

```
Residuals:
```

Min	1Q	Median	3Q	Max
-6.708	-4.134	-1.527	1.516	81.674

```
Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	9.4993	0.7304	13.006	<2e-16 ***
x	-1.5509	0.1683	-9.213	<2e-16 ***

```
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 7.965 on 504 degrees of freedom
```

```
Multiple R-squared:  0.1441, Adjusted R-squared:  0.1425
```

```
F-statistic: 84.89 on 1 and 504 DF, p-value: < 2.2e-16
```

```
> lmplot(crim,rad)
```

```
Call:
```

```
lm(formula = y ~ x, data = Boston)
```

```
Residuals:
```

Min	1Q	Median	3Q	Max
-10.164	-1.381	-0.141	0.660	76.433

```
Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-2.28716	0.44348	-5.157	3.61e-07 ***
x	0.61791	0.03433	17.998	< 2e-16 ***

```
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 6.718 on 504 degrees of freedom
```

```
Multiple R-squared:  0.3913, Adjusted R-squared:  0.39
```

```
F-statistic: 323.9 on 1 and 504 DF, p-value: < 2.2e-16
```

```
> lmplot(crim,tax)
```

```
Call:
```

```
lm(formula = y ~ x, data = Boston)
```

```
Residuals:
```

Min	1Q	Median	3Q	Max
-12.513	-2.738	-0.194	1.065	77.696

```
Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-8.528369	0.815809	-10.45	<2e-16 ***
x	0.029742	0.001847	16.10	<2e-16 ***

```
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 6.997 on 504 degrees of freedom
```

```
Multiple R-squared:  0.3396, Adjusted R-squared:  0.3383
```

```
F-statistic: 259.2 on 1 and 504 DF, p-value: < 2.2e-16
```

```
> lplot(crim,ptratio)
```

```
Call:
```

```
lm(formula = y ~ x, data = Boston)
```

```
Residuals:
```

Min	1Q	Median	3Q	Max
-7.654	-3.985	-1.912	1.825	83.353

```
Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-17.6469	3.1473	-5.607	3.40e-08 ***
x	1.1520	0.1694	6.801	2.94e-11 ***

```
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 8.24 on 504 degrees of freedom
```

```
Multiple R-squared:  0.08407, Adjusted R-squared:  0.08225
```

```
F-statistic: 46.26 on 1 and 504 DF, p-value: 2.943e-11
```

```
> lplot(crim,black)
```

```
Call:
```

```
lm(formula = y ~ x, data = Boston)
```

```
Residuals:
```

Min	1Q	Median	3Q	Max
-13.756	-2.299	-2.095	-1.296	86.822

```
Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	16.553529	1.425903	11.609	<2e-16 ***
x	-0.036280	0.003873	-9.367	<2e-16 ***

```
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 7.946 on 504 degrees of freedom
```

```
Multiple R-squared:  0.1483, Adjusted R-squared:  0.1466
```

```
F-statistic: 87.74 on 1 and 504 DF, p-value: < 2.2e-16
```

```
> lplot(crim,lstat)
```

```

Call:
lm(formula = y ~ x, data = Boston)

Residuals:
    Min       1Q   Median       3Q      Max
-13.925  -2.822  -0.664   1.079   82.862

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept) -3.33054    0.69376  -4.801 2.09e-06 ***
x             0.54880    0.04776  11.491 < 2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 7.664 on 504 degrees of freedom
Multiple R-squared:  0.2076, Adjusted R-squared:  0.206
F-statistic: 132 on 1 and 504 DF, p-value: < 2.2e-16

```

```
> lmplot(crim,medv)
```

```

Call:
lm(formula = y ~ x, data = Boston)

Residuals:
    Min       1Q   Median       3Q      Max
-9.071 -4.022 -2.343   1.298  80.957

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept) 11.79654    0.93419  12.63  <2e-16 ***
x           -0.36316    0.03839   -9.46  <2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 7.934 on 504 degrees of freedom
Multiple R-squared:  0.1508, Adjusted R-squared:  0.1491
F-statistic: 89.49 on 1 and 504 DF, p-value: < 2.2e-16

```

(15.b)

Please find attached below the .R file.



```
> summary(lmplotall)
```

Call:

```
lm(formula = crim ~ ., data = Boston)
```

Residuals:

Min	1Q	Median	3Q	Max
-9.924	-2.120	-0.353	1.019	75.051

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	17.033228	7.234903	2.354	0.018949	*
zn	0.044855	0.018734	2.394	0.017025	*
indus	-0.063855	0.083407	-0.766	0.444294	
chas	-0.749134	1.180147	-0.635	0.525867	
nox	-10.313535	5.275536	-1.955	0.051152	.
rm	0.430131	0.612830	0.702	0.483089	
age	0.001452	0.017925	0.081	0.935488	
dis	-0.987176	0.281817	-3.503	0.000502	***
rad	0.588209	0.088049	6.680	6.46e-11	***
tax	-0.003780	0.005156	-0.733	0.463793	
ptratio	-0.271081	0.186450	-1.454	0.146611	
black	-0.007538	0.003673	-2.052	0.040702	*
lstat	0.126211	0.075725	1.667	0.096208	.
medv	-0.198887	0.060516	-3.287	0.001087	**

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 6.439 on 492 degrees of freedom

Multiple R-squared: 0.454, Adjusted R-squared: 0.4396

F-statistic: 31.47 on 13 and 492 DF, p-value: < 2.2e-16

The following predictors have p-value as:

Rad – 6.46e-11

Dis – 0.000502

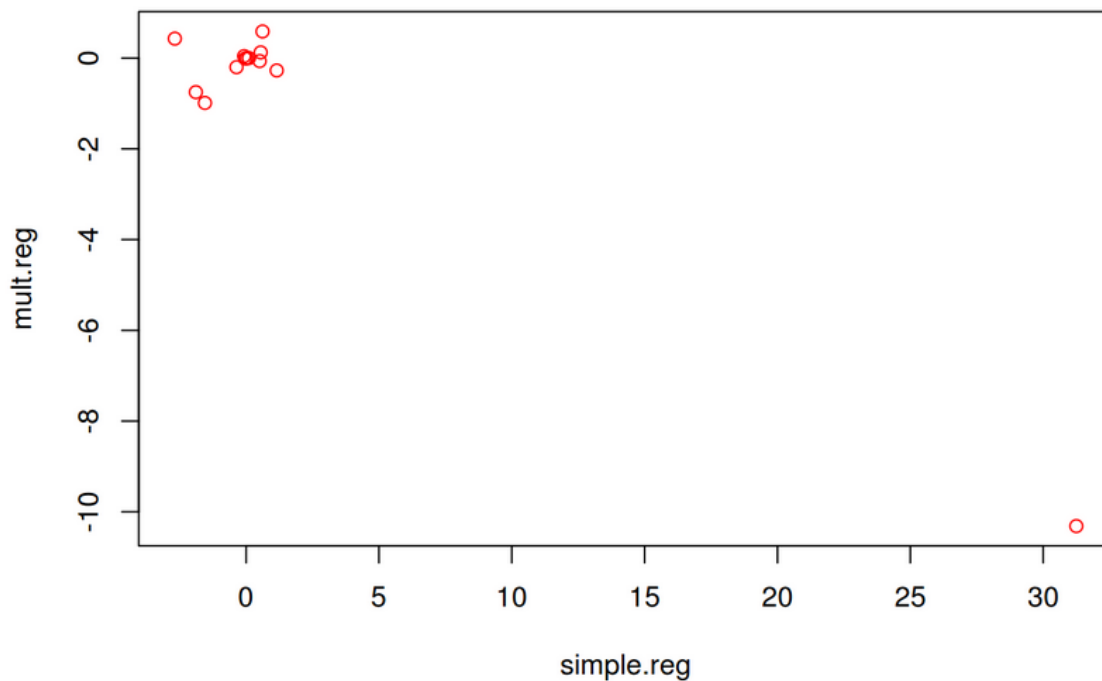
Med – 0.001087

Zn – 0.017025

Black – 0.040702

The above predictors reject the Null hypothesis because they have low p-value.

15.c)



There is a difference between the simple and multiple regression coefficients. This difference is due to the fact that in the simple regression case, the slope term represents the average effect of an increase in the predictor, ignoring other predictors. In contrast, in the multiple regression case, the slope term represents the average effect of an increase in the predictor, while holding other predictors fixed. It does make sense for the multiple regression to suggest no relationship between the response and some of the predictors while the simple linear regression implies the opposite because the correlation between the predictors show some strong relationships between some of the predictors.

```
>cor(Boston[-c(1,4)])
```

```
##           zn      indus      nox      rm      age      dis
## zn      1.0000000 -0.5338282 -0.5166037  0.3119906 -0.5695373  0.6644082
## indus  -0.5338282  1.0000000  0.7636514 -0.3916759  0.6447785 -0.7080270
## nox    -0.5166037  0.7636514  1.0000000 -0.3021882  0.7314701 -0.7692301
## rm     0.3119906 -0.3916759 -0.3021882  1.0000000 -0.2402649  0.2052462
## age    -0.5695373  0.6447785  0.7314701 -0.2402649  1.0000000 -0.7478805
## dis     0.6644082 -0.7080270 -0.7692301  0.2052462 -0.7478805  1.0000000
## rad    -0.3119478  0.5951293  0.6114406 -0.2098467  0.4560225 -0.4945879
## tax    -0.3145633  0.7207602  0.6680232 -0.2920478  0.5064556 -0.5344316
## ptratio -0.3916785  0.3832476  0.1889327 -0.3555015  0.2615150 -0.2324705
## black   0.1755203 -0.3569765 -0.3800506  0.1280686 -0.2735340  0.2915117
## lstat   -0.4129946  0.6037997  0.5908789 -0.6138083  0.6023385 -0.4969958
## medv    0.3604453 -0.4837252 -0.4273208  0.6953599 -0.3769546  0.2499287
##           rad      tax      ptratio      black      lstat      medv
## zn      -0.3119478 -0.3145633 -0.3916785  0.1755203 -0.4129946  0.3604453
## indus    0.5951293  0.7207602  0.3832476 -0.3569765  0.6037997 -0.4837252
## nox      0.6114406  0.6680232  0.1889327 -0.3800506  0.5908789 -0.4273208
## rm      -0.2098467 -0.2920478 -0.3555015  0.1280686 -0.6138083  0.6953599
## age      0.4560225  0.5064556  0.2615150 -0.2735340  0.6023385 -0.3769546
## dis     -0.4945879 -0.5344316 -0.2324705  0.2915117 -0.4969958  0.2499287
## rad      1.0000000  0.9102282  0.4647412 -0.4444128  0.4886763 -0.3816262
## tax      0.9102282  1.0000000  0.4608530 -0.4418080  0.5439934 -0.4685359
## ptratio  0.4647412  0.4608530  1.0000000 -0.1773833  0.3740443 -0.5077867
## black   -0.4444128 -0.4418080 -0.1773833  1.0000000 -0.3660869  0.3334608
## lstat    0.4886763  0.5439934  0.3740443 -0.3660869  1.0000000 -0.7376627
## medv    -0.3816262 -0.4685359 -0.5077867  0.3334608 -0.7376627  1.0000000
```

So for example, when “age” is high there is a tendency in “dis” to be low, hence in simple linear regression which only examines “crim” versus “age”, we observe that higher values of “age” are associated with higher values of “crim”, even though “age” does not actually affect “crim”. So “age” is a surrogate for “dis”; “age” gets credit for the effect of “dis” on “crim”.

15.d)

Please find attached below the .R file



isl15d.R

‘Chas’ does not support

```
> lmpoly(crim,zn)
```

Call:

```
lm(formula = y ~ poly(x, 3), data = Boston)
```

Residuals:

```
      Min      1Q Median      3Q      Max
-4.821 -4.614 -1.294  0.473 84.130
```

Coefficients:

```
      Estimate Std. Error t value Pr(>|t|)
(Intercept)   3.6135     0.3722   9.709 < 2e-16 ***
poly(x, 3)1  -38.7498     8.3722  -4.628 4.7e-06 ***
poly(x, 3)2   23.9398     8.3722   2.859 0.00442 **
poly(x, 3)3  -10.0719     8.3722  -1.203 0.22954
```

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 8.372 on 502 degrees of freedom
Multiple R-squared: 0.05824, Adjusted R-squared: 0.05261
F-statistic: 10.35 on 3 and 502 DF, p-value: 1.281e-06

```
> lmpoly(crim,indus)
```

Call:

```
lm(formula = y ~ poly(x, 3), data = Boston)
```

Residuals:

```
      Min      1Q Median      3Q      Max
-8.278 -2.514  0.054  0.764 79.713
```

Coefficients:

```
      Estimate Std. Error t value Pr(>|t|)
(Intercept)    3.614     0.330  10.950 < 2e-16 ***
poly(x, 3)1    78.591     7.423  10.587 < 2e-16 ***
poly(x, 3)2   -24.395     7.423  -3.286 0.00109 **
poly(x, 3)3   -54.130     7.423  -7.292 1.2e-12 ***
```

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 7.423 on 502 degrees of freedom
Multiple R-squared: 0.2597, Adjusted R-squared: 0.2552
F-statistic: 58.69 on 3 and 502 DF, p-value: < 2.2e-16

```
> lmpoly(crim,nox)
```

Call:

```
lm(formula = y ~ poly(x, 3), data = Boston)
```

Residuals:

```
      Min      1Q Median      3Q      Max
-9.110 -2.068 -0.255  0.739 78.302
```

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	3.6135	0.3216	11.237	< 2e-16	***
poly(x, 3)1	81.3720	7.2336	11.249	< 2e-16	***
poly(x, 3)2	-28.8286	7.2336	-3.985	7.74e-05	***
poly(x, 3)3	-60.3619	7.2336	-8.345	6.96e-16	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 7.234 on 502 degrees of freedom
Multiple R-squared: 0.297, Adjusted R-squared: 0.2928
F-statistic: 70.69 on 3 and 502 DF, p-value: < 2.2e-16

```
> lmpoly(crim,rm)
```

Call:

```
lm(formula = y ~ poly(x, 3), data = Boston)
```

Residuals:

Min	1Q	Median	3Q	Max
-18.485	-3.468	-2.221	-0.015	87.219

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	3.6135	0.3703	9.758	< 2e-16	***
poly(x, 3)1	-42.3794	8.3297	-5.088	5.13e-07	***
poly(x, 3)2	26.5768	8.3297	3.191	0.00151	**
poly(x, 3)3	-5.5103	8.3297	-0.662	0.50858	

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 8.33 on 502 degrees of freedom
Multiple R-squared: 0.06779, Adjusted R-squared: 0.06222
F-statistic: 12.17 on 3 and 502 DF, p-value: 1.067e-07

```
> lmpoly(crim,age)
```

Call:

```
lm(formula = y ~ poly(x, 3), data = Boston)
```

Residuals:

Min	1Q	Median	3Q	Max
-9.762	-2.673	-0.516	0.019	82.842

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	3.6135	0.3485	10.368	< 2e-16	***
poly(x, 3)1	68.1820	7.8397	8.697	< 2e-16	***
poly(x, 3)2	37.4845	7.8397	4.781	2.29e-06	***
poly(x, 3)3	21.3532	7.8397	2.724	0.00668	**

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 7.84 on 502 degrees of freedom
Multiple R-squared: 0.1742, Adjusted R-squared: 0.1693
F-statistic: 35.31 on 3 and 502 DF, p-value: < 2.2e-16

> lmpoly(crim,dis)

Call:

lm(formula = y ~ poly(x, 3), data = Boston)

Residuals:

Min	1Q	Median	3Q	Max
-10.757	-2.588	0.031	1.267	76.378

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	3.6135	0.3259	11.087	< 2e-16	***
poly(x, 3)1	-73.3886	7.3315	-10.010	< 2e-16	***
poly(x, 3)2	56.3730	7.3315	7.689	7.87e-14	***
poly(x, 3)3	-42.6219	7.3315	-5.814	1.09e-08	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 7.331 on 502 degrees of freedom
Multiple R-squared: 0.2778, Adjusted R-squared: 0.2735
F-statistic: 64.37 on 3 and 502 DF, p-value: < 2.2e-16

> lmpoly(crim,rad)

Call:

lm(formula = y ~ poly(x, 3), data = Boston)

Residuals:

Min	1Q	Median	3Q	Max
-10.381	-0.412	-0.269	0.179	76.217

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	3.6135	0.2971	12.164	< 2e-16	***
poly(x, 3)1	120.9074	6.6824	18.093	< 2e-16	***
poly(x, 3)2	17.4923	6.6824	2.618	0.00912	**
poly(x, 3)3	4.6985	6.6824	0.703	0.48231	

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 6.682 on 502 degrees of freedom
Multiple R-squared: 0.4, Adjusted R-squared: 0.3965
F-statistic: 111.6 on 3 and 502 DF, p-value: < 2.2e-16

```
> lmpoly(crim,tax)
```

Call:

```
lm(formula = y ~ poly(x, 3), data = Boston)
```

Residuals:

Min	1Q	Median	3Q	Max
-13.273	-1.389	0.046	0.536	76.950

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	3.6135	0.3047	11.860	< 2e-16	***
poly(x, 3)1	112.6458	6.8537	16.436	< 2e-16	***
poly(x, 3)2	32.0873	6.8537	4.682	3.67e-06	***
poly(x, 3)3	-7.9968	6.8537	-1.167	0.244	

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 6.854 on 502 degrees of freedom
Multiple R-squared: 0.3689, Adjusted R-squared: 0.3651
F-statistic: 97.8 on 3 and 502 DF, p-value: < 2.2e-16

```
> lmpoly(crim,ptratio)
```

Call:

```
lm(formula = y ~ poly(x, 3), data = Boston)
```

Residuals:

Min	1Q	Median	3Q	Max
-6.833	-4.146	-1.655	1.408	82.697

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	3.614	0.361	10.008	< 2e-16	***
poly(x, 3)1	56.045	8.122	6.901	1.57e-11	***
poly(x, 3)2	24.775	8.122	3.050	0.00241	**
poly(x, 3)3	-22.280	8.122	-2.743	0.00630	**

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 8.122 on 502 degrees of freedom
Multiple R-squared: 0.1138, Adjusted R-squared: 0.1085
F-statistic: 21.48 on 3 and 502 DF, p-value: 4.171e-13

```
> lmpoly(crim,black)
```

Call:

```
lm(formula = y ~ poly(x, 3), data = Boston)
```

Residuals:

Min	1Q	Median	3Q	Max
-13.096	-2.343	-2.128	-1.439	86.790

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	3.6135	0.3536	10.218	<2e-16	***
poly(x, 3)1	-74.4312	7.9546	-9.357	<2e-16	***
poly(x, 3)2	5.9264	7.9546	0.745	0.457	
poly(x, 3)3	-4.8346	7.9546	-0.608	0.544	

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 7.955 on 502 degrees of freedom
Multiple R-squared: 0.1498, Adjusted R-squared: 0.1448
F-statistic: 29.49 on 3 and 502 DF, p-value: < 2.2e-16

```
> lmpoly(crim,lstat)
```

Call:

```
lm(formula = y ~ poly(x, 3), data = Boston)
```

Residuals:

Min	1Q	Median	3Q	Max
-15.234	-2.151	-0.486	0.066	83.353

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	3.6135	0.3392	10.654	<2e-16	***
poly(x, 3)1	88.0697	7.6294	11.543	<2e-16	***
poly(x, 3)2	15.8882	7.6294	2.082	0.0378	*
poly(x, 3)3	-11.5740	7.6294	-1.517	0.1299	

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 7.629 on 502 degrees of freedom
Multiple R-squared: 0.2179, Adjusted R-squared: 0.2133
F-statistic: 46.63 on 3 and 502 DF, p-value: < 2.2e-16

```
> lmpoly(crim,medv)
```

Call:

```
lm(formula = y ~ poly(x, 3), data = Boston)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-24.427	-1.976	-0.437	0.439	73.655

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	3.614	0.292	12.374	< 2e-16	***
poly(x, 3)1	-75.058	6.569	-11.426	< 2e-16	***
poly(x, 3)2	88.086	6.569	13.409	< 2e-16	***
poly(x, 3)3	-48.033	6.569	-7.312	1.05e-12	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 6.569 on 502 degrees of freedom
Multiple R-squared: 0.4202, Adjusted R-squared: 0.4167
F-statistic: 121.3 on 3 and 502 DF, p-value: < 2.2e-16

The below table shows which predictors have non-linear association based on p-value.

Predictor	Linear	$B_0+B_1*x+B_2*x^2$ (Non-linear)	$B_0+B_1*x+B_2*x^2+Bx^3$ (Non-linear)
Zn	NO	YES	NO
Indus	YES	YES	YES
Nox	YES	YES	YES
Rm	YES	YES	NO
Age	YES	YES	YES
Dis	YES	YES	YES
Rad	YES	YES	NO
Tax	YES	YES	NO
PtRatio	YES	YES	YES
Black	YES	NO	NO
Lstat	YES	YES	NO
medv	YES	YES	YES